

Toward million communicating threads

Hoang-Vu Dang
Department of Computer
Science
University of Illinois at
Urbana-Champaign
hdang8@illinois.edu

Marc Snir
Department of Computer
Science
University of Illinois at
Urbana-Champaign
snir@illinois.edu

William Gropp
Department of Computer
Science
University of Illinois at
Urbana-Champaign
wgropp@illinois.edu

ABSTRACT

We present the design, implementation and analysis of an efficient runtime which supports of million concurrently communicating threads using Message Passing Interface.

1. INTRODUCTION

Message Passing Interface (MPI) has been implemented by several vendors. On Blue Gene machine, MPICH [5] is the de-factor implementation, likewise on Infiniband machine, we have MVAPICH [7]. Cray has their own customization on top of MPICH called CrayMPI; Intel similarly implements IntelMPI. OpenMPI is another effort from Open-source collaboration of several industry and academic vendors [4]. Each vendor relies on their best knowledge of the underlying computing machine and network architecture to optimize their implementation. For examples, the implementation might reduce latency by accessing directly the low-level network-API or optimize critical path using specialized instructions.

Recently, because of the increasing in the number of processor cores within a computing board, there raises a need of running MPI efficiently in threaded environments. Although MPI specification requires all of its procedures to be thread-safe, MPI implementations are not originally designed to optimize for using multiple threads: OpenMPI latest version still does not provide a stable support of `MPI_THREAD_MULTIPLE` [10], MPICH supports thread-safety via a coarse-grain locking which essentially serialized all MPI code. The research to replace coarse-grain locking is still undergoing [1], but the progress is slow because the existing software stack is only optimized for single-core performance. For example, there are significant number of global variables shared among various component of a MPICH. Global variables which implements stack, queue or hash-table can be replaced with concurrent data structure, others might be protected with finer-grains lock. However, locks and concurrent data structure has overhead could be added to the critical path and reduces performance of single-threaded ap-

plications.

On the other hand, communication in multiple threads is troublesome for many reasons. Data movement between processors could cache trashing or false sharing if taken care of. Processors can becomes under-utilize because threads execute I/O operations can be preempted. Resource management have to consider locality and efficient concurrent accesses since they are shared among processors. Last but not least, kernel thread context-switching is costly, increasingly with the number of threads. Several studies have reported significant overhead for performing MPI concurrently in many threads [6, 1].

Because of the lack of efficient multi-threaded implementation, to our best knowledge, there is no production application that uses the thread-safety feature of MPI. What has been ended up is application running MPI in a “Thread Funnelled” mode i.e. only one thread is allowed to execute MPI procedure; other threads are used only for computation. This leads to a fairly restrictive programming model, in which communication and computation code appears in different phases or procedures. This approach often uses MPI non-blocking features to allow overlapping between computation and communication and hope the MPI runtime will be able to make communication progress behind the scene. In practice, to support better asynchronous communication, user of MPI implementations often need to enable a flag to request the runtime to spawn a dedicated polling thread for progressing the communication (e.g., `MPICH_ASYNC_PROGRESS` flag in MPICH/MVAPICH).

The solution that we have seen for existing MPI implementation to adopt multi-core machine is thus ad-hoc and inefficient. In this paper, we present an approximate implementation of MPI point-to-point communication which takes into account of multi-core machine by design. Our implementation is approximate since we shall relax some semantics requirement of MPI. We argue that our design and algorithms for this relaxation is able to achieve efficient implementation, yet do not restrict applications from using MPI effectively. More complex MPI semantics and functionality can be built on top of our point-to-point communication. By using this bottom-up strategy, we provide an insight on the minimal cost to implement MPI, accurately to the number of memory accesses. We summarize our most significant contributions as follows:

- A light-weight thread scheduler using bit-vector which achieves single-instruction for signalling continuation.
- An constant time overhead algorithm for MPI point-to-point communication utilising a specially designed

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions.acm.org.

Euro MPI Aug XX-YY, 2016, XXX, YYY, ZZZ

© 2016 ACM. ISBN 123-4567-24-567/08/06.

DOI: 10.475/123.4

concurrent wait-free hash-table.

- A resource-aware locality-aware concurrent memory pool for packet management.
- A MPI runtime design for scaling up to million communicating threads.

Our paper is organized as follows. The next section describes our runtime architecture from the high-level design to individual components. We present our implementation details and optimization in Section 3. Section 4 discusses experiment and results. Section 5 gives an overview of some related works. Finally Section 6 concludes our study.

2. RUNTIME ARCHITECTURE

2.1 High-Level System Design

We first identify the complication of MPI code is on the message matching mechanism. MPI point-to-point procedures (Send/Recv) facilitate matching messages by tagging each with an integer in addition to destination and communicator. The implementation uses several queues data structure to store incoming requests and messages. For example, the unexpected queue store arrived messages without a matching request; the posted queue will store the pending requests without matching messages. Depending on different arrival order, traversing, insertion or deletion of entries into queues will be issued. In a concurrent setting, these are critical section that need to be protected. An efficient lock-free control of these operations are non-existent even though individual queue could be implemented efficiently. Moreover, when the number of concurrent communication increases, traversing a queue requires linear time to the size of pending requests.

We propose an implementation MPI based on the following assumptions:

- Large number of concurrent threads Threads are lightweight; they are scheduled by a user-space scheduler that understands synchronization objects.
- The implementation is optimized for the case where `MPLANY_SOURCE` and `MPLANY_TAG` are not used.
- No concurrent send and receive having the same signature, thus we do not yet worry about ordering.
- The NIC can be accessed in user space; it has its own routing tables, in order to translate ranks in `MPL_COMM_WORLD` to physical node addresses; it has page tables, in order to translate virtual addresses to physical addresses.
- We consider only x86_64 machine, however the technique is general enough to port to other architecture which supports atomic exchange and atomic bit manipulation.

The key data structure for communication is a lock-free hash table which stores both unsolicited incoming messages and outstanding receives. Since we assume there are no wildcard, we can hash by source and tag. We focus on critical operations, the creation of communicators or of datatypes is not (yet) address. We also ignore, for the time being, one-sided operations.

The runtime is an integrated system of a user-level thread (ULT) scheduler and a communication server. The ULT provides a low latency thread management while the server is a kernel-thread dedicating to message delivery. The two components interact using the aforementioned hash-table. Our primary goal is to design an algorithm and optimize to minimize the latency of this interaction. A failure to achieve this goal simply adds extraneous overhead to single threaded application and might not achieve significant benefit that justifies its cost. We shall also show that our design allows optimization space to be reduced to optimizing a few necessary operations.

Resource management is another important issue. Aside from communication context such as connections and NIC-provided control data which we maintains per communication server, we have to maintain a pool of packets structure. The pool is the second shared data structure of the runtime. The two basic operations are `get` and `ret` which obtains and returns a packet from/to the pool respectively. Each packet is used for storing control and message data to deliver to network device.

Figure 1 shows the overall architecture of our described runtime system.

2.2 Point-to-Point Message Passing Algorithms

We now describe an algorithm and its data structure that facilitates the point-to-point communication in Message Passing semantics. Our algorithm relies on a specialized concurrent hash-table H defined as follows.

We denote a tuple $(k, v) \in H$ when v is stored in the hash-table using the key k . At initialization, for any key k , only the tuple (k, \perp) are stored in the table. The hash-table has two operations:

- $H.\text{insert}(k, v)$: attempt to replace value of key k by v .
- $H.\text{empty}(k)$: replace value of key k by \perp

Additionally, let H_t denote a state of H which is a set of all key-value pair stored in H at time t , H_{t_0} denote the state of H before a operation and H_{t_1} denotes the state of H after a operation; \mathbb{K}, \mathbb{V} denote the key and entry space. In a sequential history, we have the following legal semantics:

$$\text{insert}(k, v) = \begin{cases} v \iff (k, \perp) \in H_{t_0}, (k, v) \in H_{t_1} \\ v' \iff (k, v') \in H_{t_0}, (k, v') \in H_{t_1} \end{cases} \quad (1)$$

$$\text{empty}(k) = \text{success} \iff (k, v) \in H_{t_0}, (k, \perp) \in H_{t_1} \quad (2)$$

$$\forall k_0 \in \mathbb{K}, \nexists (v_0, v_1) \in \mathbb{V} \mid v_0 \neq v_1 \wedge \{(k_0, v_0), (k_0, v_1)\} \subset H_t \quad (3)$$

That is, (1) means that `insert` succeeds only if the entry being stored with k at the time of insertion is \perp . In that case, (k, \perp) is replaced with (k, v) . Otherwise, the operations fail and existing value is returned i.e. no changed are made to H . In contrast, `erase` as in (2) is always successful, which replaces any value v with the input k with \perp , essentially removing v from the hash-table. The (3) equation is the consistency requirement for hash-table, which means for each key we allow only one value associated with it.

In a concurrent setting, we further require the hash-table to be *linearizable*. This means we ensure firstly safety and correctness property. Secondly, we ensure operations take

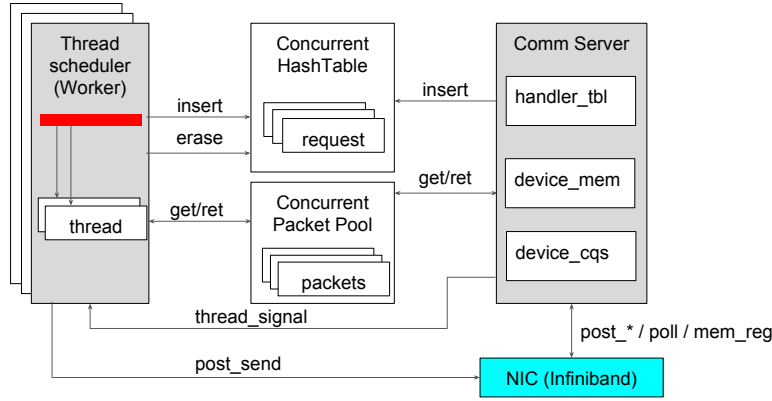


Figure 1: MPI Runtime Architecture for multi-threaded executions

affect in a real-time order. This is a strong guarantee however is necessary to implement MPI semantics which have operations executed in program order. Lastly, linearizability is *composable* which allows us to correctly use the hash-table to implement other concurrent objects.

Message delivery in general can be implemented in two way: *eager* or *rendevouz* protocol. Eager protocol is used when the target buffer is not known to the sender. Thus, an copy is made into an intermediate buffer - usually packed with other control data to deliver to the network. This also allows the Send operation to return immediately as the buffer can be reused. This protocol however becomes inefficient when message size gets larger, typically larger than the L2 or L3 cache size i.e. the cost of data movement is significant. When this is the case, we switch to *rendevouz* protocol in which the data is delivered directly from the source buffer to the target buffer by the NIC thus saving extra copies. The protocol however requires some control messages to exchange control data and signal completion.

2.2.1 Eager protocol

The pseudocode for eager protocol is listed in Algorithm 1 and 2 for worker thread and communication server respectively. There is not much for the eager send, the thread can return immediately without any waiting since the content of the data shall be copied over to the packet data for transferring. Only the receiving algorithm needs some elaborations. The basic idea here is that the thread and the communication server could use the results from hash-table insertion to coordinate the matching of messages and request.

The hashing function takes the request signature (rank, tag, communicator) and produces a 64-bit key for inserting into the table. It is easy to see that when there is a matching message and request, both the server and a thread are having the same signature and key. Depending on which comes first the algorithm goes into two cases. If the communication server succeeds with the hash insertion, the server knows the receiving request has not been posted and it returns immediately. A thread later comes, eventually fails the insertion but finds a packet with the needed data to work with. On the other hand, if the thread is the one who succeeds, the request is then inserted into the hash-table with the synchronization object. The thread is yielded by executing `ThreadWait`. The

Algorithm 1 Eager-message send/rcv for thread

```

1: procedure SEND-EAGER( $b, s, d$ )           ▷ : buffer, size,
   destination signature
2:    $p = \text{pkpool.get}()$ 
3:   Set packet header  $p.d$  to  $d$ 
4:   Copy  $b$  to  $p.b$ 
5:   Post  $p$  to network for send.
6: end procedure
7:
8: procedure RECV-EAGER( $b, s, d$ ) ▷ : buffer, size, source
   signature
9:   Create a request  $r = (b, s, d, (\omega, \Gamma))$ 
10:  Create hash-value  $v$  from  $r$ .
11:  Create hash-key  $k$  from  $d$ .
12:   $v' = H.\text{insert}(k, v)$ 
13:  if  $v' \neq v$  then                               ▷ : insertion fail
14:    Copy  $v'.p.b$  to  $b$  ▷ : message arrived, copy data.
15:     $\text{pkpool.ret}(p)$ 
16:  else                                             ▷ : insertion success
17:     $\text{ThreadWait}()$  ▷ : message not arrived, wait.
18:  end if
19:   $H.\text{erase}(k)$ 
20: end procedure

```

Algorithm 2 Eager-message packet handler for communication server

```

1: procedure RECV-EAGER-PACKET( $p$ )
2:   Create hash-value  $v$  from  $p$ .
3:   Create hash-key  $k$  from  $p.d$ .
4:    $v' = H.\text{insert}(k, v)$ 
5:   if  $v' \neq v$  then                               ▷ : insertion fail.
6:     Copy  $p.b$  to  $v'.r.b$  ▷ : thread arrived, copy data.
7:      $\text{ThreadSignal}(r.\Gamma)$ 
8:      $\text{pkpool.ret}(p)$ 
9:   else                                             ▷ : insertion success.
10:    return ▷ : thread not arrive, nothing to do.
11:  end if
12: end procedure

```

worker is now free to do other works. When the associated packet arrived, the server will fail the insertion and find the request with the attached synchronization object. It now can mark the thread as schedulable using `ThreadSignal` so when the worker is available it will pick up the works with the available data.

Since there is no need for linear searching, our matching algorithm achieves a constant time complexity as long as operation of thread scheduler, pool and hash-table is also constant time.

2.2.2 Rendezvous protocol

A rendezvous protocol have the same algorithm as short protocol after we have exchanged the control message via eager-protocol. The control messages includes two messages: a RTS (ready-to-send) issued by the sender, a RTR (ready-to-receive) issued by the receiver. By then, the sender and the receiver has known the addresses of each other buffer and they can perform communication. The data transfer could be optimized further using the Remote Direct Memory Access (RDMA) feature of modern Network Interface Controller (NIC). Several researchs has focused on optimizing this operation.

As an effect of our assumption, in this protocol we can save one control message i.e. the RTS. The reason is we do not have wild-card, the sender and receiver knows exactly their target. We only require the receiver to send its buffer to the sender. The sender can follow up by issuing an RDMA. Analogous to the eager protocol, whenever the sender or receiver is required to wait for a matching message it will perform `ThreadWait`, and later when the message has arrived the communication server will perform a `ThreadSignal`.

Specifically, the sender waits when the RTR message has not arrived or when RDMA is pending. In both situation the server wake up the sender thread when RDMA has completed. The receiver after issuing the RTR will wait until the signal that RDMA has finished arrived and data is now available.

2.3 Critical Operations Discussion

Clearly, optimizing the hash-table, the packet pool and the thread scheduler operations are the key to achieve good performance for our protocols. One could achieve $O(1)$ amortized in complexity however, an efficient implementation requires optimizing for the constant factor. Ideally, we want these operations to be *wait-free* (every operation has a bound on number of steps before completion) to ensure progress and guarantee system-wide throughput. Moreover, we want to minimize memory accesses thus reducing possible cache miss. In the next section, we describe the implementaion of each of the three components and show how we could optimize toward our goals.

2.4 Supporting wild-cards

Although it is out of the scope of this paper to deal with wild-cards, we provide here some possible options to mitigate existing applications. A MPI-complete implementation could be built on top of our foundation but possibly losing some of the properties that we want for performance. However, among the applications that we have looked at, these mitigation approaches are often sufficient.

An application depending on `MPI_ANY_TAG` can be rewritten such that the receiver can have a tag (Note that the

sender always has a valid tag). There are two cases:

- Number of expecting tags are bounded and deterministic: one can spawn a number of threads, each posts a `MPI_Recv` for a valid expecting tag and subsequently performs actions associating with that tag.
- The number of tags is un-bounded or non-deterministic: the sender can help by first sending its tag to the receiver via a pre-defined tag. Once sender and receiver agrees on the generated tag, they can proceed the communication.

The remedy for `MPI_ANY_SOURCE` is similar to `MPI_ANY_TAG` when there is a bounded and deterministic number of source ranks. It is harder in the other case since there is no way to notify the receiver of its sender unless there is a 3rd-party channel. For example, we could introduce a special `MPI_Send_anysource` procedure, in which the sender expects the receiver to post a `MPI_ANY_SOURCE`. In that case, when the sender receives this special packet, it will ignore the rank of the sender and uses a special tag associated with `MPI_ANY_TAG` to match into the hash-table.

When there is concurrent send/recv with the same signature (rank, tag, communicator), one can add to the tag value a session counter increasing atomically with each send/recv posted message to ensure order. When this is not possible, relaxing the hash-table to allow redundant entries with the same key is required. We refer the reader to an alternative algorithm introduced recently by Intel to allow matching using such hash-tables [3].

3. RUNTIME IMPLEMENTATION AND OPTIMIZATION

3.1 Concurrent Hash-Table

Our concurrent hash-table is not deviated very much from conventional concurrent hash-table. We however find opportunity to optimize further by taking advantage of some semantics requirement that we have mentioned.

Firstly, we use a spinlock per bucket. This is a viable option since we could control the table size to reduce the hash collision to minimal. The table size is related directly to how many concurrent operations which is controlled by our packet pool size. When there is no collision, the conflict can only happen between communication server and a thread when both try to insert into the same bucket at the same time. With a small number of conflicts, a spinlock is sufficient for synchronization and allow a very simple design.

Secondly, we design each bucket as an 4-entry array. Each entry consists of two 64-bit words. Within each 4-entry, one entry will be used as *control entry* and the other three are used for *data entry*. The control entry has an atomic flag for spin locking, and a pointer to point to the next 4-entry bucket in case we have more than 3 collisions. A data entry consists of two 64-bit words of key and value pair. In total, 4-entry takes 64-bytes and typically fits in a cache line. Thus we suffer only 1 cache miss when trying to lock the bucket and the data can be read without more cache misses.

Lastly, the `insert` operation also returns the address of the associated entry having the key in memory. This allows the `empty` operation to be a single instruction/single memory write to set the value associated with key to \perp . This is only possible when only one thread can write to an entry key

which is true when no concurrent communication with the same tag is allowed.

In conclusion, with the above optimization, the **empty** operation is lock-free and wait-free, costs only 1 memory write. The **insert** operation is cache-friendly and is typically wait-free, costs also 1 memory write in the common case of no hash collision. Including the spin-lock, our cost can be as low as two memory writes.

3.2 Thread scheduler

The thread scheduler in the system need to support at least the two most important operations for synchronization purpose: **ThreadWait** and **ThreadSignal**. We provide 3 different implementation of thread scheduler: POSIX Thread, Argobots, and a customized thread scheduler using bitvector.

For both POSIX Thread and Argobots scheduler, the two operations are implemented using **condition variable**. Condition variable is a generic container which is used as a building block for many other synchronization primitives. It typically requires a mutex and a queue traversal to find the waiting thread. An alternative is to use a busy-waiting synchronization flag which has lower latency, however this is considered worst for scaling since the processor spends useless time polling, thus we do not follow this path.

The different between POSIX Thread and Argobots is that the Argobots scheduler implements User-Level Thread (ULT), which maintains a fixed number of worker and schedule *light-weight thread* on top of that. The purpose is to reduce the overhead of context switching because a context-switch in ULT bypass the kernel and thus lower the latency to less than a hundred cycles.

Additionally, to further show that improving **ThreadWait** and **ThreadSignal** is critical to the algorithm, we design and implement a specialized thread scheduler based on bit-vector, namely *Fult*. The context-switching mechanism is similar to that of Argobots. However, rather than using a queue as traditional ULT - we indicates a schedulable thread by a single bit. When a worker is created, it is given a unique worker id, denotes as ω . When a thread is created by a worker, it is also assigned an unique id Γ . A pair (ω, Γ) uniquely defines a thread in the system at a point in time. During creation time, the associated Γ bit in its bitvector structure is marked so that when the worker is free it will attempt to schedule that thread to run. Algorithms 3 further describes our scheduling algorithms.

The scheduler works at 64-threads granularity instead of 1-thread traditionally. By using an atomic exchange instruction to swap the interested word to local variable, we are able to continuously perform read/write from/to this variable without accessing the main memory. This is a work-around since we could not found suitable atomic read-modify-write instructions operating at the bit level. An improvement in instruction usage could further improve this algorithm.

Algorithms 4 describes the two thread operations. **ThreadWait** simply switch back to the worker scheduler based on current worker id (i.e. ω) On the other hand, **ThreadSignal** shall access the bit-vector and atomically set the bit in the appropriate word based on the the same information stored in the synchronization object. ω and Γ are stored in thread local-storage of the worker, and is updated whenever a new context is switch to.

Note that the algorithm works because it is guaranteed

Algorithm 3 Thread scheduler

```

1: procedure SCHEDULING( $\omega$ ,  $V$ )  $\triangleright$  worker, bit-vector
2:   while ! $\omega$ .stop do  $\triangleright$  loop until user ask to stop
3:     for word in  $V$  do
4:       if word  $\neq 0$  then
5:         localWord = 0
6:         AtomicExchg(word, localWord)
7:         while localWord64 > 0 do
8:           b = FindFirstSet(localWord)
9:           localWord = FlipBit(localWord, b)
10:          ContextSwitch(b)
11:        end while
12:      end if
13:    end for
14:  end while
15: end procedure

```

Algorithm 4 Thread Operations

```

1: procedure THREADWAIT
2:   ContextSwitch( $\omega$ )
3: end procedure
4:
5: procedure THREADSIGNAL(sync)
6:   AtomicBitSet(sync. $\omega$ .V, sync. $\Gamma$ )
7: end procedure

```

that the scheduler will see the **ThreadSignal** after performing **ThreadWait**. It is errornous to perform **ThreadSignal** multiple times on the same object since they will be treated as one or more signal depending on the time the worker looks at the bit-vector. Thus, to support efficient non-blocking communication and different flavor of **MPI_Wait***, we have to build a more powerful synchronization.

Iterating over the bit-vector word by word is more expensive when there is more threads. We eliminate this problem by implement a hierachical design of bit-vector. That is, we could use an additional bit-vector as a hint to index into the original bit-vector. Specifically, each bit in the secondary bit-vector indicates which word in the first level may have a schedulable thread. More specifically, **ThreadSignal** will perform a bit set first into the first level then a bit set into the second level. The scheduler will first look into the second level to find a potential word and go directly to that word to look for schedulable threads. Using this scheme, suppose we allows s_2 words in the second level and each bit represents s_1 words in the first level, the total number of threads we can support is $s_1 \times s_2 \times 64^2$. Typically we can choose $s_1 = s_2 = 8$ since 8 words fit into a cache line. In that case, it allows 256K concurrent threads per worker.

Although fairness could be an issue, our scheduler maintain progress property. If a thread is marked as schedulable, it's eventually will be scheduled in a bounded number of steps. To show that this is true, consider the atomic exchange as taking a snapshot of the global state. In this snapshot, if a thread is marked it will be scheduled after all threads having smaller index are scheduled, which is bounded by the maximum number of allowed concurrent threads decided by the size of the bit-vectors.

3.3 Concurrent Packet Pool

In general, the packet pool can be implemented using a




Figure 2: Packet life cycle: (1) A thread sending data obtain a packet from the pool; (2) The thread writes data into the packet and submit it to the NIC; (1', 2') Communication server obtains a packet for receiving data and posts the packet to the NIC; (3) The server polls the NIC and obtains the packet back (either send/rcv packet). (4) If the packet is for sending, it is returned to the pool; If the packet is for receiving and there is a request, the server copies the data and returns the packet to the pool; (4') If the packet is for receiving and the there is no request, the packet is inserted to hash-table. (5') The thread takes packet from the hash-table copies the data before returning it back to the pool.

lockfree stack. A pool `ret` is translated to a stack `push`, and `get` is translated to a stack `pop`. At the initialization, a fixed number of packet is initialized from the main memory and `push` to the container. The Last-In-First-Out (LIFO) property allows good temporal locality for writing/reading to/from the content of a data packet. In a single-threaded environment this design is sufficient for good performance, but not for the case of multi-core/multi-threaded. Consider a packet recently being used by a thread and returned to the pool, this packet could be subsequently obtained by a different thread running in a different core. This leads to several cache misses since the cacheline is momentarily owned by the first thread. An example is when two threads running in two different cores alternatively perform `MPI_Send`.

Figure 2 explains how different components in our system might change the affinity of data inside a packet. The figure shows the life of a packet from the time it leaves the pool until it is returned. As explained in the figure, when a packet returns to the pool, it can only has the affinity of either the communication server or one of the worker thread. More specifically, if the packet is used for sending data, it shall have the locality of the thread; if the packet is used for receiving data, it shall have the locality of either the thread or the server depending on whether the associated request and the packet is received first. Moreover, when the packet is posted by the server to the NIC for receiving data, the packet shall lose any affinity it has since the NIC will perform a write over the packet.

Analysing the packet life cycle naturally motivates us a new design for the packet pool: split the centralized packet management into a private pool per worker. Initially, there is a fixed number of packets for each of the pool. At runtime,

we allow moving packets among those pools via *resource-stealing*.

The private pool consists of packets that are identified as having affinity of the associated hardware thread. It is implemented as a fixed size double-ended queue (deque). The deque has three main operations: `popTop`, `pushTop`, and `popBottom` which allows LIFO accessing at the one end and taking items from the other end. The idea is that at the bottom of the deque are packets that are least recently used and are better candidates for other purposes.

When a packet is needed, a `popTop` is performed with the private pool of the current worker, which guarantees giving the last packet that was read or written by the same worker. If the private pool is empty, It performs a *steal* by randomly choosing another private pool and perform the `popBottom` operation. Our algorithms essentially performs resource-balancing among the set of packets by moving it around pools with the locality taken into account. Currently we implement the pool using a simple fixed-size ring-buffer with a spinlock.

4. PERFORMANCE EVALUATION

4.1 Component overheads

In this section we evaluate our implementation of each individual component. Understanding them individually gives us an idea on the minimum cost of the overall system.

4.1.1 Concurrent Hash-Table

To evaluate the overhead due to hash-table operations, we measure the latency on hash-table operations in the two following scenarios when performing them in a number of POSIX threads:

- A thread performs `insert` when there is no item with the same key.
- A thread performs `insert` and there is already items in the hash-table, and it subsequently performs `erase`.

The two scenarios represent the only two possible runtime execution of the hash-table in our algorithms, thus measuring the latency in both cases give us an idea on how it will add to the overhead of MPI procedure. To justify the benefit of customization, we also compare ourselves to two popular general purposes hash-tables: libcuckoo implementing cuckoo hashing (*ch*) [8], and TBB concurrent hashmap (*tbb*) [9]. The experiment is to create threads and have each performed a fixed number of insertion (and erase) with different key to collect the mean latency. We ran the experiment for 1000 times and collect the summaries of the latency of a thread. Between each run, we also perform a cache invalidation.

Figure 3 shows the result of our experiment. In both cases, both TBB concurrent hash map and libcuckoo shows inconsistent latency when there is more concurrent threads, which is the result of conflicts. Since our hash-table is optimized for these patterns, our latency is almost always as low as 50 nanosecond. This latency is consistent with our expectation that our cost for the hash-table is equivalent to 2 memory writes.

4.1.2 Thread scheduler

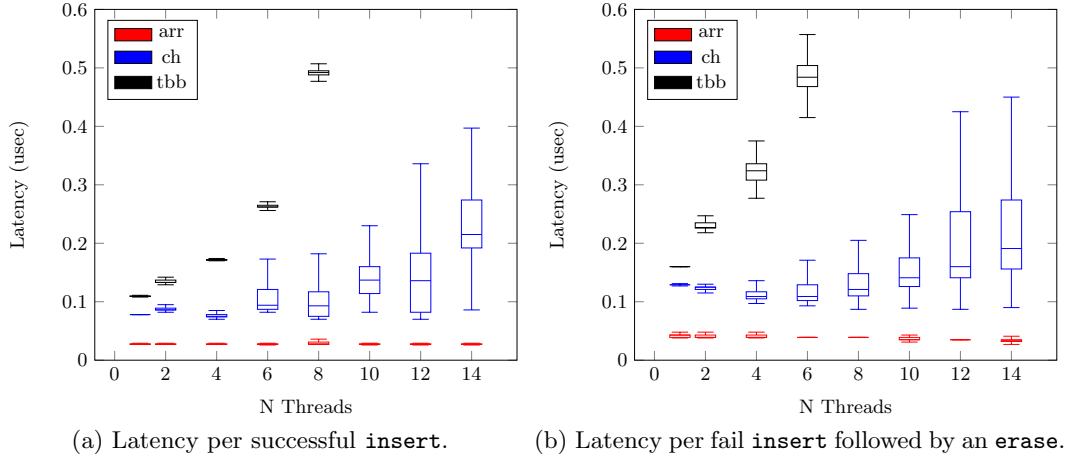


Figure 3: Latency of our hash-table implementation (*arr*) in comparison to libcuckoo (*ch*) and tbb concurrent hash map (*tbb*). Each hash-table is created with the initial size of 2^{16} , the number of insertion per thread is chosen so that there is enough room and no expansion is required. TBB is also compiled with *tbb-malloc* to improve performance. Latency exceeds an 0.5 microsecond is not shown.

	POSIX thread	Argobots	Fult
Scheduling	0.75	0.08	0.02
Signal	1.15	0.30	0.01
Total	1.90	0.38	0.03

Table 1: Break down of thread scheduler overheads, shown in *usec*.

To evaluate the thread scheduler minimum overhead, we measure separately the two operations: 1) How fast can a thread be scheduled - by performing a large number of yield within a thread and obtain the entire runtime. 2) The cost of a `ThreadSignal` by getting the runtime latency when performing operation on a thread of a different worker. Table 1 shows our results, Fult scheduler achieves the overall cost of 30 nanosecond, about $10\times$ that of Argobots and $60\times$ of POSIX threads.

4.1.3 Concurrent Packet Pool

The overhead due to packet pool is measured as the sum of the latency of `get` and `ret` operation. We evaluate this quantity by performing a random number of `get` followed by the same number of `ret` in each thread. To match better with the real workload, we also perform a random sleep in between the two group of operations. Further, The range of the random number is chosen with the same seed for each implementation and ensure threads do not request more than available packets: the maximum is number of packets divided by number of threads. We perform this experiment 1000 times and plot the summary similarly as the hash-table experiment.

The result is shown in Figure 4 in comparison with implementation using a concurrent lockfree stack and a lockfree queue. Our result for this benchmark outperforms others by a wide margin, especially when the number of threads increase. The lockfree stack is faster than the queue at single thread, however performs worst for more than two threads since there is contention at the top of the stack. The great variation in latency per operation of a centralized pool is




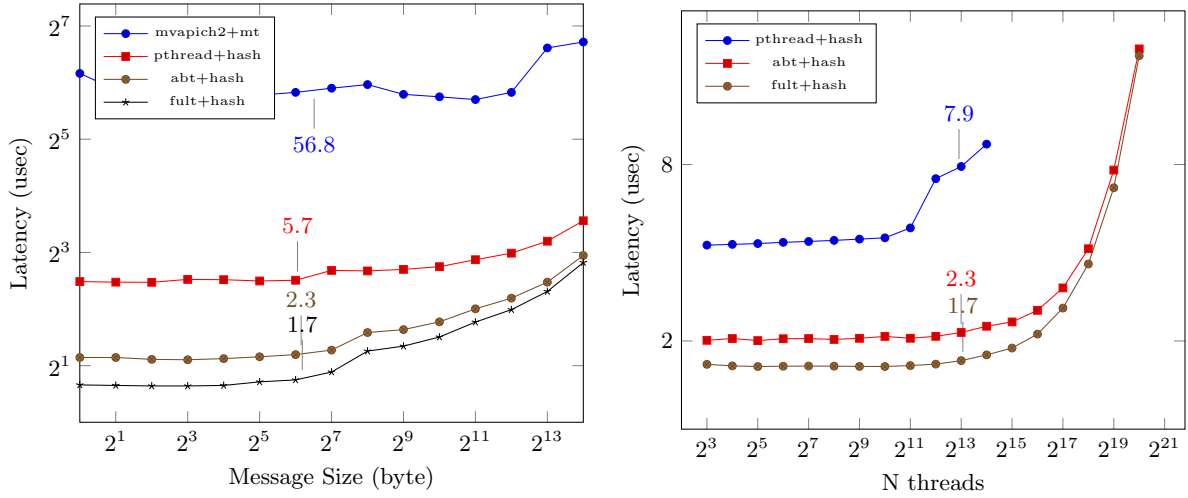
Figure 4: Latency of pool implementation vs. a lock-free pool and a lockfree queue implementation, latency higher than 1 microsecond are not shown.

clearly due to memory conflicts in this type of access patterns. Our latency is consistently in the range of 100-150 nanosecond.

4.2 Microbenchmarks and Application Evaluation

4.2.1 OSU latency benchmarks

We use OSU benchmarks [2] to evaluate the latency per `MPI_Send` or `MPI_Recv`. The single threaded test is performed using `osu_latency`, the multi-threaded test is performed with `osu_latency_mt`. We compare our best implementation (*fult+hash*) with MVAPICH2 and other possible implementation of our runtime design. Table 2 summarizes the differences of the tested runtime. *mvapich2+mt* is the MVAPICH2 with POSIX thread which is the base implementation in `osu_latency_mt`. Further, in multi-threaded



(a) Latency per message transfer for 14 threads, one per worker/socket. (b) Latency per 64-byte message transfer for upto 1M thread and round-robin assigned to worker/socket, *pthread+hash* version only works up to 16K threads due to limitation in the system.

Figure 5: Latency comparison between different MPI implementation using OSU multi-threaded latency test.

	Matching Algorithm	Scheduler
mvapich2+mt	queues	POSIX thread
pthread+hash	hash-table	POSIX thread
abt+hash	hash-table	Argobots
fult+hash	hash-table	Fult

Table 2: Summary of MPI implementation used in the evaluation, Speedup is shown for latency test in Figure 5 for message size of 64 bytes in comparison to mvapich2+mt.

test, we modify the code so that each thread uses different tags.

Performance results for multi-threaded tests are shown in Figure 5. Under tested settings, the most improvement in performance is due to the replacement of the complex queue of MPI with the hash-table. Then, the replacement of POSIX threads with a light-weight threads. Our thread-scheduler improves the latency upto 40% compared to Argobots, 3 times compared to POSIX thread scheduler. Overall, we achieve speedup of upto 45 \times compared to MVA-PICH2.

Our overheads analyzed in the previous section translate accurately to the single-threaded benchmark results in Figure 6. In general, compared to a single-threaded MPI implementation, our best implementation achieves overall overhead of 0.2 *usec* virtually tie with running MPI in a POSIX thread in MPL_THREAD_MULTIPLE mode (as in *mvapich2+mt*).

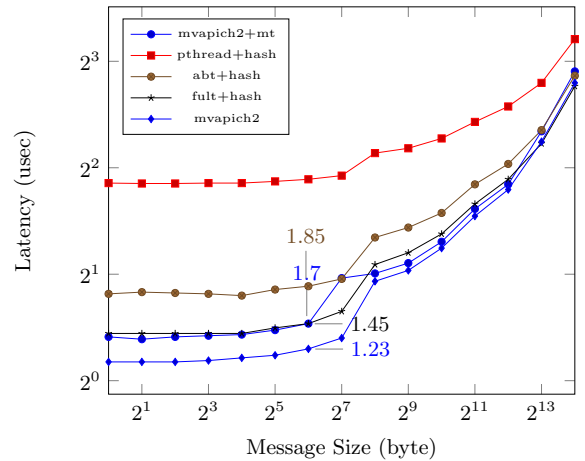


Figure 6: Latency comparison for single-threaded OSU latency test.

5. RELATED WORKS

6. CONCLUSION

7. REFERENCES

- [1] P. Balaji, D. Buntinas, D. Goodell, W. Gropp, and R. Thakur. Toward efficient support for multithreaded mpi communication. In *Recent Advances in Parallel Virtual Machine and Message Passing Interface*, pages 120–129. Springer, 2008.
- [2] P. Dhabaleswar. Osu micro-benchmarks 5.3 (03/30/16). <http://mvapich.cse.ohio-state.edu/benchmarks/>, 2016. [Online; accessed 18-April-2016].
- [3] M. Flajslik, J. Dinan, and K. D. Underwood. Mitigating MPI Message Matching Misery. In *ISC 2016*, 2016.
- [4] E. Gabriel, G. E. Fagg, G. Bosilca, T. Angskun, J. J. Dongarra, J. M. Squyres, V. Sahay, P. Kambadur, B. Barrett, A. Lumsdaine, et al. Open mpi: Goals, concept, and design of a next generation mpi implementation. In *Recent Advances in Parallel Virtual Machine and Message Passing Interface*, pages 97–104. Springer, 2004.
- [5] W. Gropp, E. Lusk, N. Doss, and A. Skjellum. A high-performance, portable implementation of the mpi message passing interface standard. *Parallel computing*, 22(6):789–828, 1996.
- [6] W. Gropp and R. Thakur. Issues in developing a thread-safe mpi implementation. In *Recent Advances in Parallel Virtual Machine and Message Passing Interface*, pages 12–21. Springer, 2006.
- [7] W. Huang, G. Santhanaraman, H.-W. Jin, Q. Gao, and D. K. Panda. Design of high performance mvapich2: Mpi2 over infiniband. In *Cluster Computing and the Grid, 2006. CCGRID 06. Sixth IEEE International Symposium on*, volume 1, pages 43–48. IEEE, 2006.
- [8] X. Li, D. G. Andersen, M. Kaminsky, and M. J. Freedman. Algorithmic improvements for fast concurrent cuckoo hashing. In *Proceedings of the Ninth European Conference on Computer Systems*, page 27. ACM, 2014.
- [9] C. Pheatt. Intel® threading building blocks. *Journal of Computing Sciences in Colleges*, 23(4):298–298, 2008.
- [10] T. O. M. Project. Is open mpi thread safe. <https://www.open-mpi.org/faq/?category=supported-systems>, 2016. [Online; accessed 9-April-2016].